

Statistics and Estimation for Computer Science



İstanbul Teknik Üniversitesi

Mustafa Kamasak, PhD



These slides are licensed under a Creative Commons Attribution 4.0 License.

License: <https://creativecommons.org/licenses/by-nc-nd/4.0/>

Version: 2022.3.15


Analysis of Categorical Data

χ^2 Statistics

- ▶ Chi-square test is used to analyze the goodness-of-fit of observed categorical data (such as gender, profession etc.) to its expected value (due to a known model, assumption etc.).
- ▶ χ^2 statistics is used to determine the goodness-of-fit:


$$\chi^2 = \sum \frac{(\text{observed value} - \text{expected value})^2}{\text{expected value}}$$

- ▶ χ^2 statistics (rv.) has χ^2 distribution. A critical value can be determined via computation or table-lookup.



- ▶ If χ^2 value for an observation is lower than the critical value, it means the observed data is consistent with the expected model/assumption. Data fits model well.
- ▶ Very small expected values in the denominator may inflate the χ^2 statistics.
- ▶ Rule of thumb: χ^2 can be used whenever the expected value for all cases (cells) are at least 5. Otherwise, cells should be combined.

χ^2 Distribution



- ▶ The shape changes with dof (k in the figure)
- ▶ In t-distribution dof changes variation, dof does not change location (always centered at 0)
- ▶ In χ^2 distribution location and variation changes with dof.

Fair Die Example

- ▶ In reality, you might never know if a die is fair or not.
- ▶ It should be decided upon observations.
- ▶ A die is fair if $P(k) = 1/6$ for $k \in \{1, \dots, 6\}$
- ▶ Assume two dice are thrown 600 times each

Die 1:

outcome (k)	1	2	3	4	5	6
observation (O_k)	102	99	105	96	104	94

Die 2:

outcome (k)	1	2	3	4	5	6
observation (O_k)	82	110	105	120	102	81

- ▶ Is die 1 fair?
- ▶ Is die 2 fair?

Fair Die Example

- ▶ In a fair die, expected number of observations for each outcome (cell) is 100.
- ▶ For die 1:

$$\chi_1^2 = \sum_{i=1} \frac{(O_i - 100)^2}{100} = 0.98$$


- ▶ For die 2:

$$\chi_2^2 = \sum_{i=1} \frac{(O_i - 100)^2}{100} = 12.14$$

- ▶ dof = # of cells - 1. Minus one is due to the fixed number of cells.
- ▶ dof = 6 - 1 = 5.

Fair Die Example – Table Lookup

Table of upper-tail area (p-value) for χ^2 distribution:



Chi-square Distribution Table

d.f.	.995	.99	.975	.95	.9	.1	.05	.025	.01
1	0.00	0.00	0.00	0.00	0.02	2.71	3.84	5.02	6.63
2	0.01	0.02	0.05	0.10	0.21	4.61	5.99	7.38	9.21
3	0.07	0.11	0.22	0.35	0.58	6.25	7.81	9.35	11.34
4	0.21	0.30	0.48	0.71	1.06	7.78	9.49	11.14	13.28
5	0.41	0.55	0.83	1.15	1.61	9.24	11.07	12.83	15.09
6	0.68	0.87	1.24	1.64	2.20	10.64	12.59	14.45	16.81
7	0.99	1.24	1.69	2.17	2.83	12.02	14.07	16.01	18.48
8	1.34	1.65	2.18	2.73	3.49	13.36	15.51	17.53	20.09

- ▶ For die 1: $\chi_1^2 = 0.98$ and $0.975 < p_1 < 0.95$
- ▶ For die 2: $\chi_2^2 = 12.14$ and $0.025 < p_2 < 0.05$
- ▶ Critical value $\chi_c^2 = 11.07$ for $\text{dof}=5$ and $\alpha = 0.05$
 - ▶ As $\chi_1^2 < \chi_c^2$ (or as $p_1 > \alpha$), the observed data is not sufficient to claim that the die 1 is not fair.
 - ▶ As $\chi_2^2 > \chi_c^2$ (or as $p_2 < \alpha$), there is evidence that the die 2 is not fair.

Fair Die Example – Python Implementation

```
from scipy.stats import chisquare

# observations
dice1 = [102, 99, 105, 96, 104, 94]
dice2 = [82, 110, 105, 120, 102, 81]
# expected values
exp = [100, 100, 100, 100, 100, 100]

chi2, p = chisquare(dice1, exp)
print('Dice 1: chi2 stat: %.2f p: %f\n'%(chi2, p))
chi2, p = chisquare(dice2, exp)
print('Dice 2: chi2 stat: %.2f p: %f\n'%(chi2, p))
```

Output

```
Dice 1: chi2 stat: 0.98 p: 0.964163
Dice 2: chi2 stat: 12.14 p: 0.032919
```

Fair Dice Example – Small sample size

- ▶ With 18 trials, test if Dice 3 is fair:

outcome (k)	1	2	3	4	5	6
observation (O_k)	4	0	4	4	2	4
expected	3	3	3	3	3	3

- ▶ $\chi^2 = 4.67$
- ▶ Rule of thumb: χ^2 can be used whenever the expected value for all cases (cells) are at least 5. Otherwise, cells should be combined.

outcome (k)	1 - 2	3 - 4	5 - 6
observation (O_k)	4	8	6
expected	6	6	6

- ▶ $\chi^2 = 1.33$

Mendel–Fisher Controversy



- ▶ On left: Gregor Mendel (1822–1884): Founder of genetics
- ▶ On right: Sir Ronald Fisher (1890–1962): Founder of statistics

Check out "A Statistical Model to Explain the Mendel–Fisher Controversy", Ana M. Pires and J. A. Branco,

<https://arxiv.org/pdf/1104.2975v1.pdf>

Mendel–Fisher Controversy



- ▶ On 1865, Mendel published a paper about his experiments on garden peas.
- ▶ He bred pea with pure yellow and green seeds. Then he bred them.
- ▶ In the first generation, all were yellow (as it is the dominant gene)
- ▶ In the second generation, out of 8023 peas:

Color	Green	Yellow
# of obs.	2001	6022
expected	2005.75	6017.25

Mendel–Fisher Controversy

Color	Green	Yellow
# of obs.	2001	6022
expected	2005.75	6017,25

- ▶ $\chi^2 = 0.01$ – Way too good fit to theory
- ▶ In 1936, Ronald Fisher published a statistical analysis of Mendel's experiment. He computed the p-value of the Mendel's observations as 0.00004.
- ▶ He concluded that Mendel's data is manipulated.
"The data of most, if not all, of the experiments have been falsified so as to agree closely with Mendel's expectations."
"Although no explanation can be expected to be satisfactory, it remains a possibility among others that Mendel was deceived by some assistant who knew too well what was expected."

Check if Data Comes from a Certain Distribution


- ▶ It is possible to check if data comes from a certain distribution
 - ▶ Continuous values are made discrete
 - ▶ Find number of data points within each bin
 - ▶ Find expected value of observations within each bin

$$Expected = N(F_X(upper) - F_X(lower))$$

where F_X is the cdf of the distribution.

- ▶ Compute χ^2 and find critical value χ_c^2 from the loop-up table with significance level of α .
- ▶ If $\chi^2 > \chi_c^2$, reject distribution.

Check for Normality



- Data come from Normal distribution $\sim \mathcal{N}(20, 25)$?

Check for Normality – Python Implementation

```
from scipy.stats import chisquare, norm, t
import matplotlib.pyplot as plt

# sample size
N = 1000
mu = 20
std = 5

# generate sample from N(20,25)
x = std*norm.rvs(size=N)+mu

h = plt.hist(x, alpha=0.5, color='r')
bins = h[1]
vals = h[0]

# find expected values
expected = []
for i in range(len(bins)-1):
    ulimit = (bins[i+1]-mu)/std
    llimit = (bins[i]-mu)/std
    ev = N*(norm.cdf(ulimit)-norm.cdf(llimit))
    expected.append(ev)

# perform chi2 test
chi2, p = chisquare(vals, expected)
print(chi2, p)
```


Check for Normality – Python Implementation

Range	observed	expected
3.99-7.19	4	4.5
7.19-10.39	14	22.1
10.39-13.60	72	72.8
13.60-16.80	164	160.9
16.80-20.00	257	239.1
20.00-23.20	252	239.1
23.20-26.41	148	160.8
26.41-29.61	62	72.7
29.61-32.81	24	22.1
32.81-36.02	3	4.5

- ▶ $\chi^2 = 8.42$ and $p = 0.493$
- ▶ No enough evidence to reject Normal distribution.

Check for Normality – Python Implementation

```
from scipy.stats import chisquare, norm, t
import matplotlib.pyplot as plt

# sample size
N = 1000
mu = 20
std = 5

# generate sample from t with dof=18, mu=20, std=5
# x = std*norm.rvs(size=N)+mu
x = std*t.rvs(df=18,size=N)+mu

h = plt.hist(x, alpha=0.5, color='r')
bins = h[1]
vals = h[0]

# find expected values
expected = []
for i in range(len(bins)-1):
    ulimit = (bins[i+1]-mu)/std
    llimit = (bins[i]-mu)/std
    ev = N*(norm.cdf(ulimit)-norm.cdf(llimit))
    expected.append(ev)

# perform chi2 test
chi2, p = chisquare(vals, expected)
print(chi2, p)
```

Check for Normality – Python Implementation

Range	observed	expected
4.81-7.85	16	6.4
7.85-10.89	27	26.7
10.89-13.94	88	78.3
13.94-16.98	146	160.1
16.98-20.02	220	228.8
20.02-23.06	241	228.3
23.06-26.10	143	159.1
26.10-29.14	79	77.4
29.14-32.19	27	26.3
32.19-35.23	13	6.2

- ▶ $\chi^2 = 27.08$ and $p = 0.00136$
- ▶ Reject Normal distribution for $\alpha = 0.05$.

Independence of Bivariate Categorical Data

- ▶ χ^2 statistics can also be used to investigate independence of bivariate categorical data.
- ▶ Assume we want to investigate the relation between *smoking during pregnancy* and *birth weight*.
- ▶ Observations are recorded as yes/no for smoking during pregnancy, underweight/normal/overweight for birth weights.
- ▶ Data looks like:
{(yes, normal), (yes, underweight), (no, normal), ...}
- ▶ From this data, two-way frequency table is formed as follows:

Observed	Underweight	Normal	Overweight	Total
Smoking	55	44	8	107
Non-smoking	35	62	12	109
Total	90	106	20	216

Independence of Bivariate Categorical Data

Observed	Underweight	Normal	Overweight	Total
Smoking	55	44	8	107
Non-smoking	35	62	12	109
Total	90	106	20	216

- ▶ This table is called *contingency table*¹.
- ▶ At the last row/column sum of the corresponding row/column is written.
- ▶ Grand total is written to the last cell.
- ▶ Hypothesis are:
 - ▶ H_0 : data at rows and columns are independent from each other
 - ▶ H_1 : data at rows and columns are **NOT** independent from each other
- ▶ χ^2 dof is (number of rows-1) \times (number of columns-1)
- ▶ For this example, dof=(2-1) \times (3-1)=2

¹contingency: a future event or circumstance which is possible but cannot be predicted with certainty.

Independence of Bivariate Categorical Data

Observed	Underweight	Normal	Overweight	Total
Smoking	55	44	8	107
Non-smoking	35	62	12	109
Total	90	106	20	216

- ▶ If H_0 is correct, how should 90 underweight births distributed into smoking and non-smoking moms?
 - ▶ $90 \times 107/216$ should go to smoking cell,
 - ▶ $90 \times 109/216$ should go to non-smoking cell,
- ▶ If H_0 is correct, how should 107 smoking moms distributed into underweight/normal/overweight cells?
 - ▶ $107 \times 90/216$ should go to underweight cell,
 - ▶ $107 \times 106/216$ should go to normal cell,
 - ▶ $107 \times 20/216$ should go to overweight cell,
- ▶ The expected value for cell at row i and column j is:

$$E_{ij} = \frac{(\text{row total } i) \times (\text{column total } j)}{\text{grand total}}$$

Independence of Bivariate Categorical Data

Exp./Obs.	Underweight	Normal	Overweight
Smoking	45.58/55	52.51/44	9.91/8
Non-smoking	45.42/35	53.49/62	10.09/12

- Note that expected values in all cells are ≥ 5 .

```
from scipy.stats import chi2_contingency,
import matplotlib.pyplot as plt
import numpy as np

expected = np.array([ [45.58, 52.51, 9.91], [45.42, 53.49, 10.09] ])
observed = np.array([ [55, 44, 8], [35, 62, 12] ])

r = chi2_contingency(observed, expected)
print('chi2 stat: %.2f p-value: %f\n'%(r[0], r[1]))
```

- $\chi^2 = 8.28$ and p-value=0.0159.
- Hence H_0 should be rejected for significance level of $\alpha = 0.05$.

Yates Correction

- ▶ Consider 2×2 contingency matrix, the test statistics

$$\chi^2 = \sum_i \frac{(O_i - E_i)^2}{E_i}$$

- ▶ $O_i \sim \text{Binomial}(N, p)$
- ▶ Binomial distr is approximated with normal ²
- ▶ Remember continuity correction when a discrete dist is approximated using cont. distr.
- ▶ For 2×2 contingency matrix, Yates correction is performed as follows ³

$$\chi^2 = \sum_i \frac{(|O_i - E_i| - 0.5)^2}{E_i}$$

²Recall lower limit of 5 for each cell

³Yates, F (1934). "Contingency table involving small numbers and the chi-sq test". Supplement to the Journal of the Royal Statistical Society 1(2): 217–235. JSTOR

Yates Correction

- ▶ Yates's correction prevents overestimation of statistical significance (underestimation of p-value) for small sample size
- ▶ Use Yate's correction when at least one cell of the table has an expected count smaller than 5
- ▶ Yates's correction is criticized for overcorrection ⁴

⁴Larntz, K. (1978). Small sample comparisons of exact levels for chi-square goodness of fit statistics. *Journal of the American Statistical Association*, 73, 253-263.

Regression

Regression Analysis

- ▶ Regression⁵ analysis analyzes relation between two or more variables

$$r(X) = E(Y|X = x) = \mu_{Y|X}$$

where

X is called regressor, input, feature, predictor etc.

Y is called response, output, prediction

$r(X)$ is called regression function

- ▶ Regression is used for
 - ▶ Analyze the relation between variables
 - ▶ To form a model \implies predict a respond corresponding to an input

⁵term introduced by Francis Galton when he realized sons of extremely long/short fathers regress to mean.

Simple Linear Regression Model

- ▶ Consider a linear relation between regressor and response

$$Y = \beta_0 + \beta_1 X + \epsilon$$

where

$\beta_0, \beta_1 \in \mathbb{R}$ are constant coefficients,

ϵ is noise (error) with $\mu_\epsilon = 0$ and σ_ϵ^2 is not a function of X .

- ▶ Typically Gaussian noise is assumed $\epsilon \sim \mathcal{N}(0, \sigma_\epsilon^2)$. Why? Multiple sources of error are superimposed \implies CLT
- ▶ The regression function is

$$\begin{aligned} r(x) &= E(Y|X = x) \\ &= E(\beta_0 + \beta_1 x + \epsilon) \\ &= \beta_0 + \beta_1 x \end{aligned}$$

Simple Linear Regression Model

- ▶ Coefficients β_0 and β_1 are the parameters of simple regression model
- ▶ They are estimated from a sample of observations

$$[(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)]$$

- ▶ Least squares estimation: Minimize the squared error

$$\mathcal{C}(x, y) = \sum_i \epsilon_i^2 = \sum_i (y_i - \beta_0 - \beta_1 x_i)^2$$

and

$$\hat{\beta}_0, \hat{\beta}_1 = \arg \min_{\beta_0, \beta_1} \mathcal{C}(x, y)$$

Coefficients of Linear Regression

- ▶ Take partial derivative wrt to coefficients and make them zero

$$\left. \frac{\partial}{\partial \beta_0} \mathcal{C} \right|_{\beta_0 = \hat{\beta}_0} = 0$$

$$\left. \frac{\partial}{\partial \beta_1} \mathcal{C} \right|_{\beta_0 = \hat{\beta}_1} = 0$$

- ▶ Estimated coefficients are

$$\hat{\beta}_1 = \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sum_i (x_i - \bar{x})^2} = \frac{S_{xy}}{S_{xx}}$$

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$$

where

$$S_{xx} = \sum_i x_i^2 - N\bar{x}^2$$

$$S_{xy} = \sum_i x_i y_i - N\bar{x} \bar{y}$$

Linear Model Output and Residual

- ▶ Using parameters of model $\hat{\beta}_0$ and $\hat{\beta}_1$ model outputs can be computed for each regressor

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$$

- ▶ Difference between real response (y_i) and corresponding model output (\hat{y}_i) is called residual

$$e_i = y_i - \hat{y}_i = y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i$$

- ▶ Do not get confused
 - ▶ ϵ is error that is inherent in the system
 - ▶ e is error in the model/fit
- ▶ Residuals are the error in the model/fit. Causes are
 - ▶ Error in the system ϵ
 - ▶ Error in the model
 - ▶ Other errors in computations/code etc.

Linear Model Output and Residual

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from scipy.stats import norm


beta0 = 3
beta1 = 2
x = np.arange(1,11)
y = beta0 + beta1*x + norm.rvs(loc=0, scale=6, size=x.size)

model = LinearRegression().fit(x.reshape((-1,1)), y)
beta0_est = model.intercept_
beta1_est = model.coef_
print('estimated beta0:', beta0_est)
print('estimated beta1:', beta1_est)
y_est = beta0_est + beta1_est*x

plt.plot(x,y, 'rx', markersize=5)
plt.plot(x,y_est, 'b-')
plt.plot(x,y_est, 'go', markersize=5)
#plt.plot(x,beta0+beta1*x, 'g--')
# draw residuals
for i in range(x.size):
    plt.plot([x[i],x[i]], [y[i], y_est[i]], 'k--')


plt.legend(['data', 'model', 'model output'])
plt.show()
```


Linear Model Output and Residual



Linear Model Output and Residual

- ▶ Analysis of residuals is very useful to
 - ▶ Investigate the validity of assumptions (linearity, zero mean error, error variance not related to input)
 - ▶ Investigate regression performance (model adequacy)



Coefficient of Determination

Few definitions:

- ▶ Amount of variation in response

$$S_{YY} = \sum_i (y_i - \bar{y})^2$$

- ▶ Amount of inherent variance of response

$$SS_E = \sum_i (y_i - \hat{y}_i)^2 = \sum_i e_i^2$$

This is the variance that cannot be described by the model

- ▶ Amount of variance described by the model

$$SS_R = \sum_i (\hat{y}_i - \bar{y})^2 = S_{YY} - SS_E$$

- ▶ Coefficient of determination

$$R^2 = \frac{SS_R}{S_{YY}} = \frac{S_{YY} - SS_E}{S_{YY}}$$



Coefficient of Determination

- ▶ Coefficient of determination



$$R^2 = \frac{SS_R}{S_{YY}} = \frac{S_{YY} - SS_E}{S_{YY}}$$

- ▶ $0 \leq R^2 \leq 1$
- ▶ R^2 represent the proportion of variance in response (output) that is explained by model
- ▶ R^2 is a measure of model adequacy
 - ▶ $R^2 = 0$ model is totally inadequate to describe output
 - ▶ $R^2 = 1$ model perfectly describes output

Coefficient of Determination



Coefficient of Determination



Correlation Coefficient

- ▶ What is the correlation between y and x ?

$$\begin{aligned} r &= \frac{\text{Cov}(x, y)}{\sigma_x \sigma_y} \\ &= \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_i (x_i - \bar{x})^2 \sum_i (y_i - \bar{y})^2}} \end{aligned}$$

- ▶ What is the relation between ρ and R^2 ?

$$\begin{aligned} r &= \frac{S_{XY}}{\sqrt{S_{XX} S_{YY}}} \\ SS_E &= \frac{S_{XX} S_{YY} - S_{XY}^2}{S_{XX}} \\ |r| &= \sqrt{1 - \frac{SS_E}{S_{YY}}} = \sqrt{R^2} \end{aligned}$$

Coefficient Coefficient

- ▶ Relation between r and R^2

$$R^2 = r^2$$

- ▶ Relation between r and $\hat{\beta}_1$

$$\hat{\beta}_1 = \frac{S_{XY}}{S_{XX}} \implies r = \hat{\beta}_1 \sqrt{\frac{S_{XX}}{S_{YY}}}$$

- ▶ Relation between R^2 and $\hat{\beta}_1$

$$R^2 = \hat{\beta}_1^2 \frac{S_{XX}}{S_{YY}}$$


Residual Analysis

- ▶ Residual values can have any variance
- ▶ Normalize to $\mathcal{N}(0, 1)$
- ▶ Standardized residuals


$$e'_i = \frac{y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_i)}{\sqrt{SS_E / (N - 2)}}$$

- ▶ It is typically expected that $e'_i \in [-2, 2]$ as $P(|z| < 1.96) = 0.05$
- ▶ If there are few $e'_i \notin [-2, 2]$, corresponding outputs can be assumed to be outliers
- ▶ If there are too many $e'_i \notin [-2, 2]$, then the model is not adequate or assumptions are incorrect
- ▶ There should be no pattern in e'_i and it should look random \rightarrow Need quantitative method for measuring randomness


Patterns in Residuals




(a)



(b)



(c)



(d)

(a) Ideal case

(b) Increasing noise variance

(c) Varying noise variance

(d) Model inadequacy

Confidence Intervals for Model Parameters

- ▶ $\hat{\beta}_0$ and $\hat{\beta}_1$ are point estimates for model parameters β_0 and β_1
- ▶ What is $1 - \alpha$ confidence intervals for estimations for model parameters?
- ▶ Variance of residuals is estimated as

$$\hat{\sigma}_e^2 = \frac{1}{N-2} \sum_i e_i^2 = \frac{SS_E}{N-2}$$

$N - 2$ dof is due to estimations of $\hat{\beta}_0$ and $\hat{\beta}_1$

- ▶ Recall

$$\hat{\beta}_1 = \frac{S_{xy}}{S_{xx}}$$

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x} = \bar{y} - \frac{S_{xy}}{S_{xx}} \bar{x}$$

Confidence Intervals for Model Parameters

- ▶ If $\epsilon \sim \mathcal{N}$, least squares estimation is equal to MLE
- ▶ Recall asymptotical normality of MLE

$$\frac{\hat{\beta}_1 - \beta_1}{\text{se}(\hat{\beta}_1)} \sim \mathcal{N}(0, 1) \quad \frac{\hat{\beta}_0 - \beta_0}{\text{se}(\hat{\beta}_0)} \sim \mathcal{N}(0, 1)$$

- ▶ Estimations are unbiased

$$E(\hat{\beta}_1) = \beta_1$$

$$E(\hat{\beta}_0) = \beta_0$$

- ▶ Standard errors (se) of estimations are

$$\text{Var}(\hat{\beta}_1) = \frac{\hat{\sigma}_e^2}{S_{XX}} \implies \text{s.e.}(\hat{\beta}_1) = \sqrt{\frac{\hat{\sigma}_e^2}{S_{XX}}}$$

$$\text{Var}(\hat{\beta}_0) = \hat{\sigma}_e^2 \left(\frac{1}{N} + \frac{\bar{X}^2}{S_{XX}} \right) \implies \text{s.e.}(\hat{\beta}_0) = \sqrt{\hat{\sigma}_e^2 \left(\frac{1}{N} + \frac{\bar{X}^2}{S_{XX}} \right)}$$

Confidence Intervals for Model Parameters

- ▶ Recall asymptotic normality of MLE

$$\frac{\hat{\beta}_1 - \beta_1}{\text{se}(\hat{\beta}_1)} \sim \mathcal{N}(0, 1) \quad \frac{\hat{\beta}_0 - \beta_0}{\text{se}(\hat{\beta}_0)} \sim \mathcal{N}(0, 1)$$

- ▶ Hence $1 - \alpha$ confidence interval is

- ▶ if $N \geq 20$

$$[\hat{\beta}_i - z_{\alpha/2} \text{se}(\hat{\beta}_i), \hat{\beta}_i + z_{\alpha/2} \text{se}(\hat{\beta}_i)]$$

- ▶ if $N < 20$

$$[\hat{\beta}_i - t_{N-2, \alpha/2} \text{se}(\hat{\beta}_i), \hat{\beta}_i + t_{N-2, \alpha/2} \text{se}(\hat{\beta}_i)]$$

Hypothesis Testing for Slope

- ▶ Consider following hypothesis:

$$H_0 : \beta_1 = \beta_{1,0}$$

$$H_1 : \beta_1 \neq \beta_{1,0}$$

- ▶ Recall, if σ_e^2 is known

$$\frac{\hat{\beta}_1 - \beta_{1,0}}{\sqrt{\sigma_e^2 / S_{XX}}} \sim \mathcal{N}(0, 1)$$

- ▶ Use $\hat{\sigma}_e^2$ as estimate of σ_e^2

- ▶ Test statistics

$$T = \sqrt{\frac{(N-2)S_{XX}}{SS_E}} (\hat{\beta}_1 - \beta_{1,0})$$

- ▶ Reject H_0 is $|t_{obs}| \geq t_{N-2, \alpha/2}$ for significance level of α

Hypothesis Testing for Null Slope

- ▶ An important hypothesis is $\beta_1 = 0$ which means there is no linear relation between input/output

$$Y = \beta_0 + \beta_1 X$$

- ▶ Hypothesis is then

$$H_0 : \beta_1 = 0$$

$$H_1 : \beta_1 \neq 0$$

- ▶ Recall, if σ_e^2 is known

$$\frac{\hat{\beta}_1}{\sqrt{\sigma_e^2 / S_{XX}}} \sim \mathcal{N}(0, 1)$$

- ▶ Use $\hat{\sigma}_e^2$ as estimate of σ_e^2
- ▶ Test statistics

$$T = \sqrt{\frac{(N-2)S_{XX}}{SS_E}} (\hat{\beta}_1)$$

- ▶ Reject H_0 is $|T| \geq t_{N-2, \alpha/2}$ for significance level of α

Logistic Regression

- ▶ When response is binary ie. $y_i \in \{0, 1\}$, logistic regression is used
- ▶ Why not just use linear regression?
 - ▶ Error ϵ is not normal distributed

$$\epsilon_i = 1 - (\beta_0 - \beta_1 x_i) \text{ if } y_i = 1$$

$$\epsilon_i = -(\beta_0 - \beta_1 x_i) \text{ if } y_i = 0$$

- ▶ With linear regression, response may lie outside $[0,1]$
- ▶ Model should be nonlinear

Logit Function

- ▶ Likelihood ratio is called odds


$$\text{odds} = \frac{E(y_i)}{1 - E(y_i)} = \frac{p}{1 - p}$$

where $E(y_i) = 1 \times P(y_i = 1) + 0 \times P(y_i = 0) = p$

- ▶ Logarithm of the odds maps a probability $p \in [0, 1]$ to $[-\infty, \infty]$
- ▶ **Logistic unit** – logit – function

$$\text{logit} = \log \frac{p}{1 - p}$$


Logit Function



- To map input $x_i \in [-\infty, \infty]$ to binary response $y_i \in \{0, 1\}$, inverse logit function is required

$$\text{logit}^{-1}(x) = \text{logistic}(x)$$

Logistic Regression



- S-shaped Logistic/sigmoid function

$$\begin{aligned} P(y_i = 1) &= p \\ &= \text{logistic}(x_i) = \frac{1}{1 + \exp\{-(\beta_0 + \beta_1 x_i)\}} \end{aligned}$$

Logistic Regression

- ▶ Logistic/sigmoid function

$$\begin{aligned} P(y_i = 1) &= p \\ &= \text{logistic}(x_i) = \frac{1}{1 + \exp\{-(\beta_0 + \beta_1 x_i)\}} \end{aligned}$$

- ▶ Remember

$$\begin{aligned} \text{logit} &= \log\left(\frac{p}{1-p}\right) \\ &= \log(\exp(\beta_0 + \beta_1 x_i)) \\ &= \beta_0 + \beta_1 x_i \end{aligned}$$

Logistic Regression

- ▶ Model parameters are estimated by MLE
- ▶ No closed form expression: Use iterative methods
- ▶ Logistic regression can be used for binary classification


Logistic Regression – Example

- ▶ A group of 20 students spends between 0 and 6 hours studying for an exam.⁶
- ▶ How does the number of hours spent studying affect the probability of the student passing the exam?
- ▶ Data

```
x=[  
  (0.50,0), (0.75,0), (1.00,0), (1.25,0), (1.50,0),  
  (1.75,0), (1.75,1), (2.00,0), (2.25,1) (2.50,0),  
  (2.75,1), (3.00,0), (3.25,1), (3.50,0), (4.00,1),  
  (4.25,1), (4.50,1), (4.75,1), (5.00,1), (5.50,1)  
]
```

⁶Example from https://en.wikipedia.org/wiki/Logistic_regression

Logistic Regression – Example



- ▶ Model parameters are $\hat{\beta}_0 = -4.0777$ and $\hat{\beta}_1 = 1.5046$
- ▶ Logistic function

$$y_i = \frac{1}{1 + \exp\{-(-4.0777 + 1.5046x_i)\}}$$

Logistic Regression – Example

- ▶ Logistic function

$$y_i = \frac{1}{1 + \exp\{-(-4.0777 + 1.5046x_i)\}}$$

- ▶ Probability of passing exam with a 2-hour study is

$$y_i = \frac{1}{1 + \exp\{-(-4.0777 + 1.5046 \times 2)\}} = 0.26$$

- ▶ Probability of passing exam with a 4-hour study is

$$y_i = \frac{1}{1 + \exp\{-(-4.0777 + 1.5046 \times 4)\}} = 0.87$$

Logistic Regression – Code

```
import matplotlib.pyplot as plt
import numpy as np
from sklearn.linear_model import LogisticRegression

# number of hours studied from exam
x = np.array([0.50, 0.75, 1.00, 1.25, 1.50, \
              1.75, 1.75, 2.00, 2.25, 2.50, \
              2.75, 3.00, 3.25, 3.50, 4.00, \
              4.25, 4.50, 4.75, 5.00, 5.50])

# pass(1)/fail(0)
y = np.array([0,0,0,0,0,0,1,0,1,0,1,0,1,0,1,1,1,1,1,1])

model = LogisticRegression()
model.fit(x.reshape(-1,1),y)

print('beta0:', model.intercept_)
print('beta1:', model.coef_)
```

Multiple Linear Regression

- ▶ Typically regressor is a vector

$$Y = \beta_0 + \beta_1 X_1 + \cdots + \beta_k X_k + \epsilon$$

- ▶ Assume n responses

$$y_1 = \beta_0 + \beta_1 x_{1,1} + \cdots + \beta_k x_{1,k} + \epsilon_1$$

...

$$y_n = \beta_0 + \beta_1 x_{n,1} + \cdots + \beta_k x_{n,k} + \epsilon_n$$

- ▶ Use matrix notation

$$\mathbf{Y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}$$

where $\mathbf{Y} \in \mathbb{R}^n$, $\mathbf{X} \in \mathbb{R}^{n \times (k+1)}$, $\boldsymbol{\beta} \in \mathbb{R}^{(k+1)}$ and $\boldsymbol{\epsilon} \in \mathbb{R}^n$

Multiple Linear Regression

- ▶ In matrix notation

$$\mathbf{Y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}$$

- ▶ Parameter vector $\boldsymbol{\beta}$ is

$$\boldsymbol{\beta} = [\beta_0, \beta_1, \dots, \beta_k]^T \in \mathbb{R}^{(k+1)}$$

- ▶ Input matrix \mathbf{X} is

$$\mathbf{X} = \begin{bmatrix} 1 & x_{1,1} & x_{1,2} & \cdots & x_{1,k} \\ 1 & x_{2,1} & x_{2,2} & \cdots & x_{2,k} \\ 1 & \cdots & \cdots & \cdots & \cdots \\ 1 & x_{n,1} & x_{n,2} & \cdots & x_{n,k} \end{bmatrix} \in \mathbb{R}^{n \times (k+1)}$$

- ▶ Error vector $\boldsymbol{\epsilon}$ is

$$\boldsymbol{\epsilon} = [\epsilon_1, \epsilon_2, \dots, \epsilon_n]^T \in \mathbb{R}^n$$

Multiple Linear Regression - Least Squares Estimate

- ▶ Least squares estimate is

$$\hat{\beta} = \arg \min \mathcal{C}(\mathbf{X}, \mathbf{y})$$

where

$$\mathcal{C}(\mathbf{X}, \mathbf{y}) = (\mathbf{y} - \mathbf{X}\beta)^T (\mathbf{y} - \mathbf{X}\beta)$$

- ▶ Take partial derivative

$$\frac{\partial \mathcal{C}(\mathbf{X}, \mathbf{y})}{\partial \beta} = 0$$

- ▶ Least squares estimate is

$$\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- ▶ Prediction becomes

$$\hat{\mathbf{y}} = \mathbf{X}\hat{\beta}$$

- ▶ Residuals are

$$\mathbf{e} = \mathbf{y} - \hat{\mathbf{y}}$$

Multiple Linear Regression

- ▶ The variance estimate is

$$\sigma_e^2 = \frac{SS_E}{N - k - 1}$$

where $SS_E = e^T e$

- ▶ For $k + 1$ model parameters, residual degrees of freedom is $N - k - 1$
- ▶ Least square estimation bias and precision are

$$\begin{aligned} E(\hat{\beta}) &= \beta \\ \text{Cov}(\hat{\beta}) &= \sigma_e^2 (\mathbf{X}^T \mathbf{X})^{-1} \\ &= \sigma_e^2 \mathbf{C} \end{aligned}$$

- ▶ Standard error for $\hat{\beta}_j$ is

$$\text{se}(\hat{\beta}_j) = \sqrt{\sigma_e^2 \mathbf{C}_{jj}}$$

Polynomial Regression

- ▶ A specific form of multiple linear regression is called polynomial regression
- ▶ Preferred with curvilinear relation between regressor and response
- ▶ Regressors of a polynomial model are generated from a single input as

$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \cdots + \beta_k x^k$$

- ▶ This is a polynomial model of degree k
- ▶ Polynomial regression with degree $k = 1$ is equivalent to simple linear regression

Better Fitting with More Parameters

- ▶ As the number of inputs (parameters) in a model increases, a better fit can be obtained
- ▶ Let model 1 has K parameters and model 2 has M parameters where $K < M$
- ▶ For the same response, let \hat{y}_1 be the output of model 1 and \hat{y}_2 be the output of model 2
- ▶ Sum of squared residuals for model 2 will be smaller than model 1

$$\|y - \hat{y}_1\|_2^2 > \|y - \hat{y}_2\|_2^2$$


- ▶ Coefficient of determination of model 2 will be larger

$$R_1^2 > R_2^2$$

- ▶ With N data points, a model with $N - 1$ inputs (hence N parameters) will perfectly fit ie. $SS_E = 0$ and $R^2 = 1$

Overfitting Example

- ▶ Assume an approximately linear relation between hours of study before final exam and final exam grade
- ▶ For 10 students following observations are made



- ▶ A linear model fits well with $R^2 = 0.96$

Overfit Example

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.linear_model import LinearRegression

np.random.seed(1)

x = np.ndarray(shape=(10,10), dtype=float)
x[:,0] = np.arange(3, 13) + np.random.normal(0, 3, size=10)

for i in range(1,9):
    x[:,i] = np.random.normal(-3+i, 5, size=10)

y = x[:,0]*6 + 13 + np.random.normal(0, 10, size=10)

model = LinearRegression()
model.fit(x[:,0].reshape(-1,1), y)
hat_y = model.predict(x[:,0].reshape(-1,1))
r2 = r2_score(y, hat_y)
mse = mean_squared_error(y, hat_y)
print('degree: %d mse:%f r square: %f'%(1,mse,r2))

for i in range(1,9):
    model.fit(x[:,0:i+1], y)
    hat_y = model.predict(x[:,0:i+1])
    r2 = r2_score(y, hat_y)
    mse = mean_squared_error(y, hat_y)
    print('degree: %d mse:%f r square: %f'%(i+1,mse,r2))
```

Overfit Example

output

degree: 1	mse:31.316861	r square: 0.958046
degree: 2	mse:23.079755	r square: 0.969081
degree: 3	mse:22.268891	r square: 0.970167
degree: 4	mse:22.237426	r square: 0.970210
degree: 5	mse:20.539075	r square: 0.972485
degree: 6	mse:11.652108	r square: 0.984390
degree: 7	mse:11.581630	r square: 0.984485
degree: 8	mse:8.566025	r square: 0.988525
degree: 9	mse:0.000000	r square: 1.000000

Overfitting Example

- ▶ Even with totally random input, it is possible to increase R^2
- ▶ R^2 is a good metric for comparing models with same complexity
- ▶ R^2 is not a good metric for comparing models with different complexity \implies Use adjusted R^2
- ▶ Few data points and high model complexity may cause overfitting \implies Dropping coefficients in deep learning models
- ▶ Model memorizes data and cannot generalize
- ▶ How to decide model complexity?

Adjusted R^2

- ▶ Use a revised R^2 metric that penalizes model complexity

$$\begin{aligned} R_{adjusted}^2 &= 1 - \frac{SS_E / (N - K - 1)}{S_{YY} / (N - 1)} \\ &= 1 - \frac{(1 - R^2)(N - 1)}{N - K - 1} \end{aligned}$$

where K is the number of regressors (parameters of the model excluding β_0)

- ▶ If R^2 does not increase substantially by added complexity, adjusted R^2 will decrease

Overfit Example Revizited

output

```
degree: 1 mse:31.316861 r2: 0.958046 adj r2: 0.952802
degree: 2 mse:23.079755 r2: 0.969081 adj r2: 0.960247
degree: 3 mse:22.268891 r2: 0.970167 adj r2: 0.955251
degree: 4 mse:22.237426 r2: 0.970210 adj r2: 0.946377
degree: 5 mse:20.539075 r2: 0.972485 adj r2: 0.938091
degree: 6 mse:11.652108 r2: 0.984390 adj r2: 0.953171
degree: 7 mse:11.581630 r2: 0.984485 adj r2: 0.930181
degree: 8 mse:8.566025 r2: 0.988525 adj r2: 0.896721
degree: 9 mse:0.000000 r2: 1.000000 adj r2: nan
```

Underfitting

- ▶ Model is too simple (ie. less than required number of parameters)
→ Underfitting
- ▶ Patterns can be seen at residuals

Model Selection

- ▶ How to select a Model?
- ▶ Simpler models are preferred as they have improved generalization and lower risk of overfitting
 - ▶ Too simple \implies underfitting (model inadequacy)
 - ▶ Too complicated \implies overfitting
- ▶ If there are K inputs, there are 2^K possible models
- ▶ Two step approach for model selection
 - ▶ Define a cost function $\mathcal{C}(M_k)$ where \mathcal{C} is the cost function and M_k is the model with k regressors
 - ▶ Cost should have term for SS_E and model complexity
 - ▶ Search through models to minimize cost
- ▶ Possible cost functions are
 - ▶ Mallows's C_p statistic
 - ▶ AIC: Akaike information criteria
 - ▶ BIC: Bayesian information criteria
 - ▶ Negative R_{adj}^2 can be used as cost function

Mallow's C_p Statistic

- ▶ Mallow's C_p statistics is defined as follow

$$C(M_k) = \frac{1}{N}(SS_E + 2k\hat{\sigma}_\epsilon^2)$$

or

$$C(M_k) = \frac{SS_E}{\hat{\sigma}_\epsilon^2} - N + 2k$$

where k is the number of regressors and $\hat{\sigma}_\epsilon^2$ is the estimate of noise variance

- ▶ Two definitions will give different cost values but they will lead to same model

AIC: Akaike Information Criteria

- ▶ AIC considers model fits with MLE
- ▶ AIC cost function is

$$AIC(M_k) = N \log\left(\frac{SS_E}{N}\right) + 2k$$

- ▶ If noise has normal distribution ie $\epsilon \sim \mathcal{N}(0, \sigma_\epsilon^2)$ then Mallows C_p and AIC will lead to the same model as MLE will be equivalent to LSE

BIC: Bayesian Information Criteria

- ▶ BIC has more severe penalty for model complexity
- ▶ BIC cost function is

$$BIC(M_k) = N \log\left(\frac{SS_E}{N}\right) + k \log(N)$$

Noise Variance Estimation

- ▶ Noise variance estimate is used for mentioned cost functions
- ▶ Noise variance estimate

$$\hat{\sigma}_{\epsilon}^2 = \frac{SS_E}{N - k - 1}$$

where SS_E is obtained by all using all regressors (lowest possible SS_E with most complicated model)

Model Input Selection

- ▶ Which regressor combinations should be used for SS_E for cost computation?
- ▶ For example to compute cost of M_2 which two regressor should be used?
- ▶ Exhaustive search $C(K, k)$ possible combinations and maximize R^2
- ▶ Use stepwise regressor selection
 - ▶ forward selection
 - ▶ Start with no regressors
 - ▶ Add one regressor that increases the cost the least
 - ▶ Iterate
 - ▶ backward elimination
 - ▶ Start with all regressors
 - ▶ Eliminate one regressor that decreases the cost the most
 - ▶ Iterate
- ▶ Stepwise regression selection does not guarantee best regressor combination

Advantages and Disadvantages of Linear Regression

- ▶ Advantages:
 - ▶ Results are highly interpretable
 - ▶ Low computation cost
- ▶ Disadvantages:
 - ▶ Performance is typically worse compared to
 - ▶ Support vector regression
 - ▶ Random forest regression
 - ▶ ...⁷
 - ▶ Other issues (described at next slide)

⁷Taught in machine learning course.

Problems with Linear Regression

- ▶ Data is typically nonlinear
- ▶ Error variances typically change with data (homoscedasticity/heteroscedasticity) → Use weighted least squares (WLS)
- ▶ Very small changes in data can result in very different models → Use regularization on model parameters
- ▶ Be cautious about interpretation of model. Models show associations, not causality.

Other Types of Least Squares

- ▶ Ordinary least squares (OLS)
- ▶ Nonlinear least squares (NLS)
- ▶ Total least squares (TLS)
- ▶ Weighted least squares (WLS)
- ▶ Iterated reweighted least squares (IRLS)
- ▶ Others ⁸

⁸Check out https://en.wikipedia.org/wiki/Regression_analysis

Regression with Regularization

- ▶ Ordinary: No regression

$$\hat{\beta} = \arg \min \|\mathbf{y} - \mathbf{X}\beta\|^2$$

where

$$\|\mathbf{M}\|^2 = \mathbf{M}^T \mathbf{M}$$

- ▶ Ridge regression (L_2 regression)

$$\hat{\beta}_r = \|\mathbf{y} - \mathbf{X}\beta\|^2 + \lambda \|\beta\|^2$$

- ▶ Lasso regression (L_1 regression)

$$\hat{\beta}_l = \|\mathbf{y} - \mathbf{X}\beta\|^2 + \lambda \|\beta\|$$

where

$$\|\mathbf{M}\| = \sum_i |M_i|$$

Regression with Regularization

- ▶ Elasticnet regression (linear mixture of L_1 and L_2)

$$\hat{\beta}_I = \arg \min_{\beta} \|\mathbf{y} - \mathbf{X}\beta\|^2 + \lambda(\alpha\|\beta\| + (1 - \alpha)\|\beta\|^2)$$